

Data toolkit

May 27, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1. Demonstrate three different methods for creating identical 2D arrays in NumPy. Provide the code for each method and the final output after each method.

```
[2]: # Method 1
array_zeros = np.zeros((3,2)) # method 1 - using np.zeros
print(array_zeros)

# method 2
array_one = np.ones((3,2)) # method 2 - using np.ones
print(array_one)

# method 3
value = 5
array_full = np.full((3,2),value) # method 3 - using np.full
print(array_full)
```

```
[[0. 0.]
 [0. 0.]
 [0. 0.]]
[[1. 1.]
 [1. 1.]
 [1. 1.]]
[[5 5]
 [5 5]
 [5 5]]
```

2. Using the Numpy function, generate an array of 100 evenly spaced numbers between 1 to 10 and Reshape that 1D array into a 2D array.

```
[3]: evenly_spaced = np.linspace(1,10,100)
```

```
[4]: evenly_spaced
```

```
[4]: array([ 1.          , 1.09090909, 1.18181818, 1.27272727, 1.36363636,
          1.45454545, 1.54545455, 1.63636364, 1.72727273, 1.81818182,
          1.90909091, 2.          , 2.09090909, 2.18181818, 2.27272727,
          2.36363636, 2.45454545, 2.54545455, 2.63636364, 2.72727273,
          2.81818182, 2.90909091, 3.          , 3.09090909, 3.18181818,
          3.27272727, 3.36363636, 3.45454545, 3.54545455, 3.63636364,
          3.72727273, 3.81818182, 3.90909091, 4.          , 4.09090909,
          4.18181818, 4.27272727, 4.36363636, 4.45454545, 4.54545455,
          4.63636364, 4.72727273, 4.81818182, 4.90909091, 5.          ,
          5.09090909, 5.18181818, 5.27272727, 5.36363636, 5.45454545,
          5.54545455, 5.63636364, 5.72727273, 5.81818182, 5.90909091,
          6.          , 6.09090909, 6.18181818, 6.27272727, 6.36363636,
          6.45454545, 6.54545455, 6.63636364, 6.72727273, 6.81818182,
          6.90909091, 7.          , 7.09090909, 7.18181818, 7.27272727,
          7.36363636, 7.45454545, 7.54545455, 7.63636364, 7.72727273,
          7.81818182, 7.90909091, 8.          , 8.09090909, 8.18181818,
          8.27272727, 8.36363636, 8.45454545, 8.54545455, 8.63636364,
          8.72727273, 8.81818182, 8.90909091, 9.          , 9.09090909,
          9.18181818, 9.27272727, 9.36363636, 9.45454545, 9.54545455,
          9.63636364, 9.72727273, 9.81818182, 9.90909091, 10.          ])
```

```
[5]: reshaped_array = evenly_spaced.reshape(10,10)
```

```
[6]: reshaped_array
```

```
[6]: array([[ 1.          , 1.09090909, 1.18181818, 1.27272727, 1.36363636,
          1.45454545, 1.54545455, 1.63636364, 1.72727273, 1.81818182],
          [ 1.90909091, 2.          , 2.09090909, 2.18181818, 2.27272727,
          2.36363636, 2.45454545, 2.54545455, 2.63636364, 2.72727273],
          [ 2.81818182, 2.90909091, 3.          , 3.09090909, 3.18181818,
          3.27272727, 3.36363636, 3.45454545, 3.54545455, 3.63636364],
          [ 3.72727273, 3.81818182, 3.90909091, 4.          , 4.09090909,
          4.18181818, 4.27272727, 4.36363636, 4.45454545, 4.54545455],
          [ 4.63636364, 4.72727273, 4.81818182, 4.90909091, 5.          ,
          5.09090909, 5.18181818, 5.27272727, 5.36363636, 5.45454545],
          [ 5.54545455, 5.63636364, 5.72727273, 5.81818182, 5.90909091,
          6.          , 6.09090909, 6.18181818, 6.27272727, 6.36363636],
          [ 6.45454545, 6.54545455, 6.63636364, 6.72727273, 6.81818182,
          6.90909091, 7.          , 7.09090909, 7.18181818, 7.27272727],
          [ 7.36363636, 7.45454545, 7.54545455, 7.63636364, 7.72727273,
          7.81818182, 7.90909091, 8.          , 8.09090909, 8.18181818],
          [ 8.27272727, 8.36363636, 8.45454545, 8.54545455, 8.63636364,
          8.72727273, 8.81818182, 8.90909091, 9.          , 9.09090909],
          [ 9.18181818, 9.27272727, 9.36363636, 9.45454545, 9.54545455,
          9.63636364, 9.72727273, 9.81818182, 9.90909091, 10.          ]])
```

3. Explain the following terms:

- The difference in np.array, np.asarray and np.asanyarray.
- The difference between Deep copy and shallow copy.

ANS - Difference in np.array, np.asarray, and np.asanyarray:

- np.array

np.array always creates a new array, regardless of the input type.

It converts input data (lists, tuples, etc.) into an ndarray.

```
[7]: list_data = [1, 2, 3]

array_from_list = np.array(list_data)
```

```
[8]: array_from_list
```

```
[8]: array([1, 2, 3])
```

- np.asarray() np.asarray converts the input into an array if it's not already an array. If the input is already an array, np.asarray doesn't create a new copy; it returns the original array.

```
[9]: list_data = [1, 2, 3]

array_from_list = np.asarray(list_data)
```

```
[10]: array_from_list
```

```
[10]: array([1, 2, 3])
```

```
[11]: list_data = (1, 2, 3)

array_from_list = np.asarray(list_data)
```

```
[12]: array_from_list
```

```
[12]: array([1, 2, 3])
```

- np.asanyarray

np.asanyarray converts the input into an array if it's not already an array. Unlike np.asarray, it may not always create a new copy, even if the input is an array. It only creates a new copy if necessary to satisfy the requirements.

```
[13]: list_data = [1, 2, 3]

array_from_list = np.asanyarray(list_data)
```

```
[14]: array_from_list
```

```
[14]: array([1, 2, 3])
```

```
[15]: list_data = (1, 2, 3)
      array_from_list = np.asarray(list_data)
```

```
[16]: array_from_list
```

```
[16]: array([1, 2, 3])
```

- The difference between Deep copy and shallow copy:
- Shallow copy

Shallow copy creates a new object, but it doesn't create copies of nested objects. Instead, it copies references to the nested objects. So, changes made to the nested objects in one copy will affect the other copy.

```
[17]: import copy
      list1 = [1, [2, 3], 4]
      shallow_copy_list1 = copy.copy(list1)  # shallow copy
```

```
[18]: shallow_copy_list1
```

```
[18]: [1, [2, 3], 4]
```

- Deep copy

Deep copy creates a new object and recursively copies all nested objects within it. It means it creates an entirely new copy of the original data structure, including all nested objects. Hence, changes made in one copy will not affect the other copy.

```
[19]: import copy
      list1 = [1, [2, 3], 4]
      deep_copy_list1 = copy.deepcopy(list1)  # Deep copy
```

```
[20]: deep_copy_list1
```

```
[20]: [1, [2, 3], 4]
```

4 . Generate a 3x3 array with random floating-point numbers between 5 and 20 then, round each number in the array to 2 decimal places.

```
[21]: # Generate random floating-point numbers between 5 and 20
      random_array = np.random.uniform(5, 20, size=(3, 3))

      # Round each number to 2 decimal places
      rounded_array = np.round(random_array, decimals=2)

      print(rounded_array)
```

```
[[12.87  8.12 19.5 ]
 [16.46  7.51  8.84]]
```

```
[16.6  11.72 16.26]]
```

5. Create a Numpy array with random integers Between 1 and 10 of shape (5,6). After creating the array perform the following operations:

a) Extract all even integers from array.

b) Extract all odd integers from array.

```
[22]: random_array = np.random.randint(1,10,size = (5,6))
```

```
# Extract all even integers from the array  
even_int = random_array[random_array % 2 == 0]
```

```
# Extract all odd integers from the array  
  
odd_int = random_array[random_array % 2 !=0]
```

```
print(even_int)  
print(odd_int)
```

```
[4 2 8 2 8 6 6 2 8 6]
```

```
[7 7 1 5 5 9 1 3 5 9 7 7 5 1 7 5 3 3 5 1]
```

6 . Create a 3D NumPy array of shape (3, 3, 3) containing random integers Between 1 and 10 . Perform the following operations:

a) Find the indices of the maximum values along each depth level (third axis).

b) Perform element-wise multiplication of between both array.

```
[23]: random_3d_array = np.random.randint(1,10 , size = (3,3,3))
```

```
# a) Find the indices of the maximum values along each depth level (third axis)  
max_indices = np.argmax(random_3d_array, axis=2)
```

```
# b) Perform element-wise multiplication between the original array and its  
↪transpose  
elementwise_multiplication = random_3d_array * random_3d_array.transpose(1, 2,  
↪0)
```

```
[24]: elementwise_multiplication
```

```
[24]: array([[[25, 16, 49],  
             [16, 15, 63],  
             [42, 10, 42]],  
  
           [[ 4, 45,  6],  
            [27, 16,  4],  
            [27,  5, 64]]],
```

```
[[42, 21, 54],
 [15, 20, 24],
 [63, 24, 9]]])
```

```
[25]: max_indices
```

```
[25]: array([[1, 2, 2],
          [1, 0, 2],
          [0, 2, 0]])
```

7. Clean and transform the 'Phone' column in the sample dataset to remove non-numeric characters and convert it to a numeric data type. Also display the table attributes and data types of each column.

```
[26]: df = pd.read_csv("People Data.csv")
```

```
[27]: df.head()
```

```
[27]:
```

	Index	User Id	First Name	Last Name	Gender	\
0	1	8717bbf45cCDbEe	Shelia	Mahoney	Male	
1	2	3d5AD30A4cD38ed	Jo	Rivers	Female	
2	3	810Ce0F276Badec	Sheryl	Lowery	Female	
3	4	BF2a889C00f0cE1	Whitney	Hooper	Male	
4	5	9afFEafAe1CBBB9	Lindsey	Rice	Female	

		Email	Phone	Date of birth	\
0		pwarner@example.org	857.139.8239	27-01-2014	
1	fergusonkatherine@example.net		NaN	26-07-1931	
2	fhoward@example.org		(599)782-0605	25-11-2013	
3	zjohnston@example.com		NaN	17-11-2012	
4	elin@example.net		(390)417-1635x3010	15-04-1923	

		Job Title	Salary
0		Probation officer	90000
1		Dancer	80000
2		Copy	50000
3	Counselling psychologist		65000
4	Biomedical engineer		100000

```
[28]: df.duplicated().sum()
```

```
[28]: 0
```

```
[29]: df.dtypes
```

```
[29]: Index          int64
      User Id      object
      First Name   object
```

```
Last Name      object
Gender         object
Email          object
Phone          object
Date of birth  object
Job Title      object
Salary         int64
dtype: object
```

```
[30]: # Step 2: Remove non-numeric characters
df['Phone'] = df['Phone'].str.replace(r'\D', '', regex=True)
```

```
[31]: df['Phone']
```

```
[31]: 0      8571398239
      1           NaN
      2      5997820605
      3           NaN
      4      39041716353010
      ...
      995      0217752933
      996      0011497107799721
      997      1750774412833265
      998      9152922254
      999      079752542467259
      Name: Phone, Length: 1000, dtype: object
```

```
[32]: # Step 3: Convert to numeric data type
df['Phone'] = pd.to_numeric(df['Phone'], errors='coerce')
```

```
[33]: df['Phone']
```

```
[33]: 0      8.571398e+09
      1           NaN
      2      5.997821e+09
      3           NaN
      4      3.904172e+13
      ...
      995      2.177529e+08
      996      1.149711e+13
      997      1.750774e+15
      998      9.152922e+09
      999      7.975254e+13
      Name: Phone, Length: 1000, dtype: float64
```

```
[34]: df['Phone']
```

```
[34]: 0      8.571398e+09
      1           NaN
      2      5.997821e+09
      3           NaN
      4      3.904172e+13
      ...
      995    2.177529e+08
      996    1.149711e+13
      997    1.750774e+15
      998    9.152922e+09
      999    7.975254e+13
      Name: Phone, Length: 1000, dtype: float64
```

```
[35]: # Step 4: Display table attributes and data types
      print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Index           1000 non-null   int64
1   User Id         1000 non-null   object
2   First Name      1000 non-null   object
3   Last Name       1000 non-null   object
4   Gender          1000 non-null   object
5   Email           1000 non-null   object
6   Phone           979 non-null    float64
7   Date of birth   1000 non-null   object
8   Job Title       1000 non-null   object
9   Salary          1000 non-null   int64
dtypes: float64(1), int64(2), object(7)
memory usage: 78.2+ KB
None
```

8 . Perform the following tasks using people dataset:

- Read the 'data.csv' file using pandas, skipping the first 50 rows.
- Only read the columns: 'Last Name', 'Gender', 'Email', 'Phone' and 'Salary' from the file.
- Display the first 10 rows of the filtered dataset.
- Extract the 'Salary' column as a Series and display its last 5 values.

```
[36]: # Read the 'data.csv' file using pandas, skipping the first 50 rows.
      pd.read_csv("People Data.csv", skiprows=50)
```

```
[36]:      50  afF3018e9cdd1dA  George  Mercer  Female  \
      0  51  CccE5DAb6E288e5      Jo  Zavala  Male
```


1	52	DfBDc3621D4bcec	Joshua	Carey	Female
2	53	f55b0A249f5E44D	Rickey	Hobbs	Female
3	54	Ed71DcfaBFd0beE	Robyn	Reilly	Male
4	55	FDaFD0c3f5387EC	Christina	Conrad	Male
..
945	996	fedF4c7Fd9e7cFa	Kurt	Bryant	Female
946	997	ECddaFEDdEc4FAB	Donna	Barry	Female
947	998	2adde51d8B8979E	Cathy	Mckinney	Female
948	999	Fb2FE369D1E171A	Jermaine	Phelps	Male
949	1000	8b756f6231DDC6e	Lee	Tran	Female

	douglascontreras@example.net	+1-326-669-0118x4341	11-09-1941	\
0	pamela64@example.net	001-859-448-9935x54536	23-11-1992	
1	dianashepherd@example.net	001-274-739-8470x814	07-01-1915	
2	ingramtiffany@example.org	241.179.9509x498	01-07-1910	
3	carriecrawford@example.org	207.797.8345x6177	27-07-1982	
4	fuentesclaudia@example.net	001-599-042-7428x143	06-01-1998	
..
945	lyonsdaisy@example.net	021.775.2933	05-01-1959	
946	dariusbryan@example.com	001-149-710-7799x721	06-10-2001	
947	georgechan@example.org	+1-750-774-4128x33265	13-05-1918	
948	wanda04@example.net	(915)292-2254	31-08-1971	
949	deannablack@example.org	079.752.5424x67259	24-01-1947	

	Human resources officer	70000
0	Nurse, adult	80000
1	Seismic interpreter	70000
2	Barrister	60000
3	Engineer, structural	100000
4	Producer, radio	50000
..
945	Personnel officer	90000
946	Education administrator	50000
947	Commercial/residential surveyor	60000
948	Ambulance person	100000
949	Nurse, learning disability	90000

[950 rows x 10 columns]

```
[37]: df = pd.read_csv('People Data.csv').head()
df
```

```
[37]:
```

	Index	User Id	First Name	Last Name	Gender	\
0	1	8717bbf45cCDbEe	Shelia	Mahoney	Male	
1	2	3d5AD30A4cD38ed	Jo	Rivers	Female	
2	3	810Ce0F276Badec	Sheryl	Lowery	Female	
3	4	BF2a889C00f0cE1	Whitney	Hooper	Male	

```
4      5 9affEafAe1CBBB9    Lindsey      Rice  Female
```

```

                                Email                Phone Date of birth \
0                pwarner@example.org      857.139.8239    27-01-2014
1  fergusonkatherine@example.net                NaN      26-07-1931
2                fhoward@example.org      (599)782-0605    25-11-2013
3                zjohnston@example.com                NaN      17-11-2012
4                elin@example.net  (390)417-1635x3010    15-04-1923
```

```

                                Job Title  Salary
0      Probation officer      90000
1                Dancer      80000
2                Copy      50000
3  Counselling psychologist      65000
4      Biomedical engineer  100000
```

```
[ ]:
```

```
[38]: # Only read the columns: 'Last Name', 'Gender', 'Email', 'Phone' and 'Salary'
      ↪from the file.
df = pd.read_csv('People Data.csv', usecols=['Last Name', 'Gender', 'Email',
      ↪'Phone', 'Salary'])
df
```

```
[38]:
      Last Name  Gender                Email                Phone \
0      Mahoney    Male      pwarner@example.org      857.139.8239
1      Rivers   Female  fergusonkatherine@example.net                NaN
2      Lowery   Female      fhoward@example.org      (599)782-0605
3      Hooper    Male      zjohnston@example.com                NaN
4      Rice     Female      elin@example.net      (390)417-1635x3010
..      ...      ...
995    Bryant   Female      lyonsdaisy@example.net      021.775.2933
996    Barry    Female      dariusbryan@example.com      001-149-710-7799x721
997    Mckinney   Female      georgechan@example.org      +1-750-774-4128x33265
998    Phelps    Male      wanda04@example.net      (915)292-2254
999    Tran      Female      deannablack@example.org      079.752.5424x67259

      Salary
0      90000
1      80000
2      50000
3      65000
4     100000
..      ...
995    90000
996    50000
997    60000
```

```
998 100000
999 90000
```

```
[1000 rows x 5 columns]
```

```
[39]: # Display the first 10 rows of the filtered dataset.
df.head(10)
```

```
[39]:   Last Name  Gender      Email      Phone  Salary
0   Mahoney   Male  pwarner@example.org  857.139.8239  90000
1    Rivers  Female fergusonkatherine@example.net      NaN  80000
2    Lowery  Female  fhoward@example.org  (599)782-0605  50000
3    Hooper   Male  zjohnston@example.com      NaN  65000
4     Rice  Female  elin@example.net  (390)417-1635x3010 100000
5  Caldwell   Male  kaitlin13@example.net  8537800927  50000
6   Hoffman   Male  jeffharvey@example.com  093.655.7480x7895 60000
7  Andersen   Male  alicia33@example.org  4709522945  65000
8     Mays    Male  jake50@example.com  013.820.4758  50000
9  Mitchell   Male  lanechristina@example.net  (560)903-5068x4985 50000
```

```
[40]: # Extract the 'Salary' column as a Series and display its last 5 value
df['Salary'].tail()
```

```
[40]: 995    90000
996    50000
997    60000
998    100000
999    90000
Name: Salary, dtype: int64
```

9. Filter and select rows from the People_Dataset, where the “Last Name” column contains the name ‘Duke’, ‘Gender’ column contains the word Female and ‘Salary’ should be less than 85000 .

```
[41]: df = pd.read_csv('People Data.csv' , usecols = ['Last Name','Gender','Salary'])
df
```

```
[41]:   Last Name  Gender  Salary
0   Mahoney   Male    90000
1    Rivers  Female    80000
2    Lowery  Female    50000
3    Hooper   Male    65000
4     Rice  Female   100000
..      ...     ...     ...
995   Bryant  Female    90000
996    Barry  Female    50000
```

```

997 Mckinney Female 60000
998 Phelps Male 100000
999 Tran Female 90000

```

[1000 rows x 3 columns]

```

[42]: filterd_df = df[(df['Last Name'] == 'Duke') & (df['Gender'] == 'Female') &
    ↪ (df['Salary'] < 85000)]
filterd_df

```

```

[42]:      Last Name  Gender  Salary
45      Duke  Female  60000
210     Duke  Female  50000
457     Duke  Female  50000
729     Duke  Female  70000

```

10 . Create a 7*5 Dataframe in Pandas using a series generated from 35 random integers Between 1 to 6.

```

[43]: # Generate a series of 35 random integers between 1 and 6
random_series = pd.Series(np.random.randint(1,7 , size = 35))

# Reshape the series into a 7x5 DataFrame
df = random_series.values.reshape(7,5)

# Convert the numpy array back to a pandas DataFrame

df = pd.DataFrame(df , columns = ['A','B','C','D','E'])

# Display The DataFrame
print(df)

```

```

      A  B  C  D  E
0  1  3  3  1  5
1  5  5  3  6  3
2  5  5  4  5  3
3  4  6  2  5  1
4  2  5  2  6  1
5  6  2  5  2  6
6  1  6  3  2  4

```

11. Create two different Series, each of length 50, with the following criteria:

- The first Series should contain random numbers ranging from 10 to 50.
- The second Series should contain random numbers ranging from 100 to 1000.
- Create a DataFrame by 'jining these Series by column, and, change the names of the columns to 'col1', 'col2'.

```
[44]: # Generate the first Series with random numbers ranging from 10 to 50
series1 = pd.Series(np.random.randint(10,50,size = 50))

# Generate the second Series with random numbers ranging from 100 to 1000
series2 = pd.Series(np.random.randint(100,1000 , size = 50))

# Create a DataFrame by joining these Series by column
df = pd.DataFrame({'col1':series1, 'col2':series2})

# Display the Data Frame
print(df)
```

	col1	col2
0	26	238
1	17	548
2	46	830
3	18	924
4	16	932
5	32	732
6	47	713
7	46	389
8	44	576
9	49	518
10	35	260
11	44	479
12	16	177
13	18	592
14	23	705
15	14	539
16	25	299
17	37	955
18	10	788
19	13	539
20	43	488
21	42	371
22	12	645
23	40	595
24	15	224
25	17	124
26	12	457
27	31	483
28	29	776
29	27	587
30	31	886
31	19	668
32	36	575

33	21	721
34	36	413
35	49	439
36	19	463
37	39	209
38	16	539
39	31	529
40	48	772
41	12	396
42	16	859
43	11	481
44	25	290
45	48	414
46	33	702
47	48	650
48	39	756
49	12	831

12 . Perform the following operations using people data set:

- Delete the 'Email', 'Phone', and 'Date of birth' columns from the dataset.
- Delete the rows containing any missing values.
- Print the final output also.

```
[45]: # Read the data source
df = pd.read_csv('People Data.csv')

# Delete the 'Email', 'Phone', and 'Date of birth' columns from the dataset.
df = df.drop(columns = ['Email', 'Phone', 'Date of birth'])

# Delete the rows containing any missing values.

df = df.dropna()

# Print the final output also.

df
```

```
[45]:
```

	Index	User Id	First Name	Last Name	Gender	\
0	1	8717bbf45cCDbEe	Shelia	Mahoney	Male	
1	2	3d5AD30A4cD38ed	Jo	Rivers	Female	
2	3	810Ce0F276Badec	Sheryl	Lowery	Female	
3	4	BF2a889C00f0cE1	Whitney	Hooper	Male	
4	5	9afFEafAe1CBBB9	Lindsey	Rice	Female	
..	
995	996	fedF4c7Fd9e7cFa	Kurt	Bryant	Female	
996	997	ECddaFEDdEc4FAB	Donna	Barry	Female	

997	998	2adde51d8B8979E	Cathy	Mckinney	Female
998	999	Fb2FE369D1E171A	Jermaine	Phelps	Male
999	1000	8b756f6231DDC6e	Lee	Tran	Female

	Job Title	Salary
0	Probation officer	90000
1	Dancer	80000
2	Copy	50000
3	Counselling psychologist	65000
4	Biomedical engineer	100000
..
995	Personnel officer	90000
996	Education administrator	50000
997	Commercial/residential surveyor	60000
998	Ambulance person	100000
999	Nurse, learning disability	90000

[1000 rows x 7 columns]

13 . Create two NumPy arrays, x and y, each containing 100 random float values between 0 and 1. Perform the following tasks using Matplotlib and NumPy:

- Create a scatter plot using x and y, setting the color of the points to red and the marker style to 'o'.
- Add a horizontal line at $y = 0.5$ using a dashed line style and label it as ' $y = 0.5$ '.
- Add a vertical line at $x = 0.5$ using a dotted line style and label it as ' $x = 0.5$ '.
- Label the x-axis as 'X-axis' and the y-axis as 'Y-axis'.
- Set the title of the plot as 'Advanced Scatter Plot of Random Values'.
- Display a legend for the scatter plot, the horizontal line, and the vertical line.

```
[46]: # Generate two NumPy arrays, x and y, each containing 100 random float values
      ↪ between 0 and 1
x = np.random.rand(100)
y = np.random.rand(100)

# a) Create a scatter plot using x and y, setting the color of the points to
      ↪ red and the marker style to 'o'
plt.scatter(x, y, color='red', marker='o', label='Scatter points')

# b) Add a horizontal line at y = 0.5 using a dashed line style and label it as
      ↪ 'y = 0.5'
plt.axhline(y=0.5, color='blue', linestyle='--', label='y = 0.5')

# c) Add a vertical line at x = 0.5 using a dotted line style and label it as
      ↪ 'x = 0.5'
```

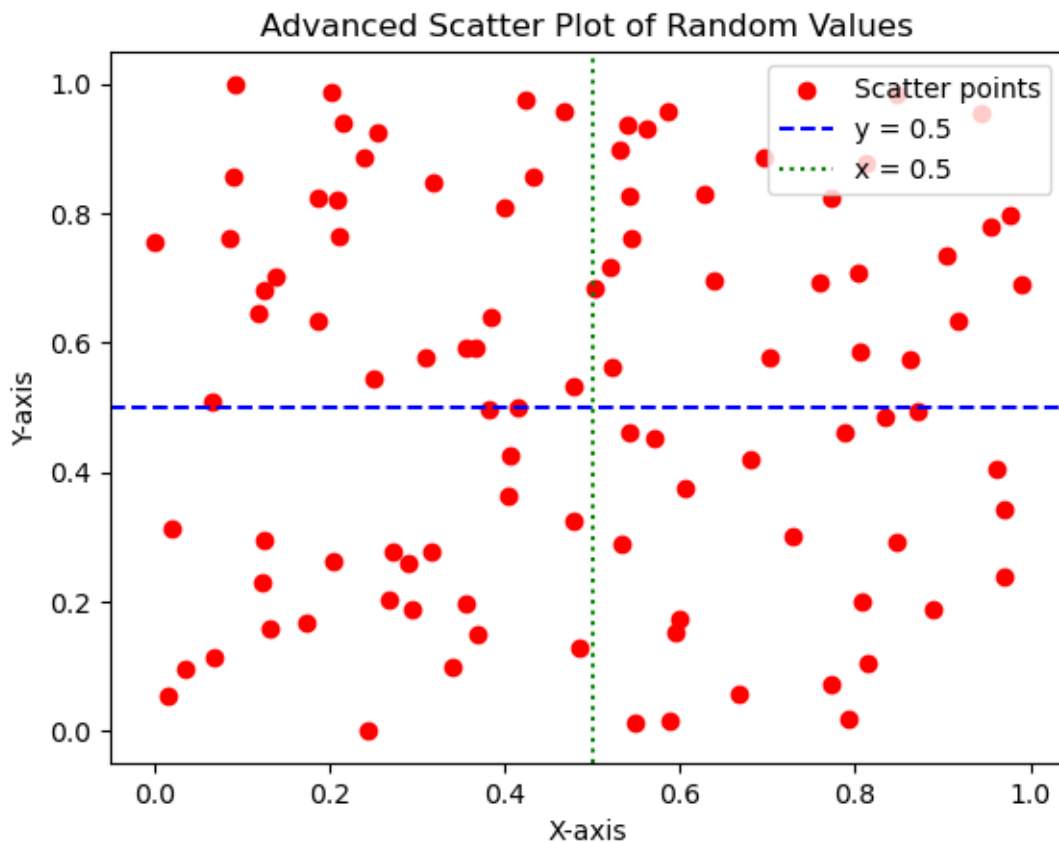
```
plt.axvline(x=0.5, color='green', linestyle=':', label='x = 0.5')

# d) Label the x-axis as 'X-axis' and the y-axis as 'Y-axis'
plt.xlabel('X-axis')
plt.ylabel('Y-axis')

# e) Set the title of the plot as 'Advanced Scatter Plot of Random Values'
plt.title('Advanced Scatter Plot of Random Values')

# f) Display a legend for the scatter plot, the horizontal line, and the
    ↪ vertical line
plt.legend()

# Show the plot
plt.show()
```



14 . Create a time-series dataset in a Pandas DataFrame with columns: 'Date', 'Temperature', 'Humidity' and Perform the following tasks using Matplotlib:

- Plot the 'Temperature' and 'Humidity' on the same plot with different y-axes (left y-axis for 'Temperature' and right y-axis for 'Humidity').

- b) Label the x-axis as 'Date'.
- c) Set the title of the plot as 'Temperature and Humidity Over Time'.

```
[47]: # Create a time-series dataset in a Pandas DataFrame
data = {
    'Date': pd.date_range(start='2024-01-01', periods=365),
    'Temperature': pd.Series(range(365)) * 2,
    'Humidity': pd.Series(range(365)) * 3
}
df = pd.DataFrame(data)

# Display Data Frame
df
```

```
[47]:
```

	Date	Temperature	Humidity
0	2024-01-01	0	0
1	2024-01-02	2	3
2	2024-01-03	4	6
3	2024-01-04	6	9
4	2024-01-05	8	12
..
360	2024-12-26	720	1080
361	2024-12-27	722	1083
362	2024-12-28	724	1086
363	2024-12-29	726	1089
364	2024-12-30	728	1092

[365 rows x 3 columns]

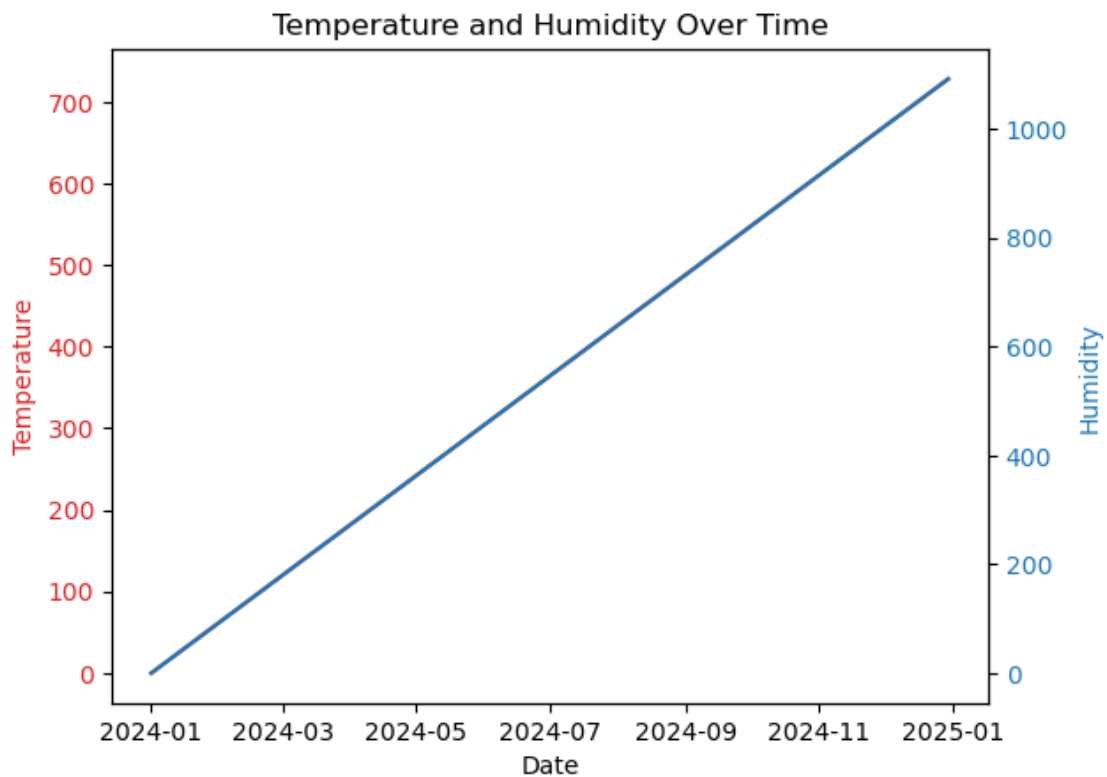
```
[48]: # Plot the 'Temperature' and 'Humidity' on the same plot with different y-axes
fig, ax1 = plt.subplots()

color = 'tab:red'
ax1.set_xlabel('Date')
ax1.set_ylabel('Temperature', color=color)
ax1.plot(df['Date'], df['Temperature'], color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()
color = 'tab:blue'
ax2.set_ylabel('Humidity', color=color)
ax2.plot(df['Date'], df['Humidity'], color=color)
ax2.tick_params(axis='y', labelcolor=color)

# Set the title of the plot
plt.title('Temperature and Humidity Over Time')
```

```
# Show the plot
plt.show()
```



15. Create a Numpy array data containing 1000 samples from a normal distribution. Perform the following tasks using Matplotlib:
 - a) Plot a histogram of the data with 30 bins.
 - b) Overlay a line plot representing the normal distribution's probability density function (PDF).
 - c) Label the x-axis as 'Value' and the y-axis as 'Frequency/Probability'.
 - d) Set the title of the plot as 'Histogram with PDF Overlay'.

```
[49]: # Generate a Numpy array containing 1000 samples from a normal distribution
data = np.random.normal(loc=0, scale=1, size=1000)

# Plot a histogram of the data with 30 bins
plt.hist(data, bins=30, density=True, alpha=0.6, color='g', label='Histogram')

# Overlay a line plot representing the normal distribution's probability
  ↳ density function (PDF)
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
```

```

p = np.exp(-0.5 * x**2) / np.sqrt(2 * np.pi)
plt.plot(x, p, 'k', linewidth=2, label='PDF')


# Label the x-axis as 'Value' and the y-axis as 'Frequency/Probability'
plt.xlabel('Value')
plt.ylabel('Frequency/Probability')

# Set the title of the plot
plt.title('Histogram with PDF Overlay')

# Display legend
plt.legend()

# Show the plot
plt.show()

```



17 . Create a Seaborn scatter plot of two random arrays, color points based on their position relative to the origin (quadrants), add a legend, label the axes, and set the title as 'Quadrant-wise Scatter Plot'.

```
[56]: # Generate random data
np.random.seed(0)
x = np.random.randn(100)
y = np.random.randn(100)

# Determine the quadrant for each point
quadrants = np.zeros_like(x)
quadrants[(x >= 0) & (y >= 0)] = 1 # Quadrant I
quadrants[(x < 0) & (y >= 0)] = 2 # Quadrant II
quadrants[(x < 0) & (y < 0)] = 3 # Quadrant III
quadrants[(x >= 0) & (y < 0)] = 4 # Quadrant IV

# Create a DataFrame for plotting
data = {'x': x, 'y': y, 'quadrant': quadrants}
df = pd.DataFrame(data)

# Set style
sns.set(style="whitegrid")


# Create the scatter plot
sns.scatterplot(x='x', y='y', hue='quadrant', palette='Dark2', data=df)

# Add legend
plt.legend(title='Quadrant')

# Label the axes
plt.xlabel('X-axis')
plt.ylabel('Y-axis')

# Set the title
plt.title('Quadrant-wise Scatter Plot')

# Show plot
plt.show()
```



18. With Bokeh, plot a line chart of a sine wave function, add grid lines, label the axes, and set the title as 'Sine Wave Function'.

```
[59]: from bokeh.plotting import figure, show
from bokeh.io import output_notebook

# Define the sine wave function
def sine_wave(x):
    return np.sin(x)

# Generate x values
x = np.linspace(0, 4*np.pi, 100)

# Generate y values using the sine wave function
y = sine_wave(x)

# Initialize the Bokeh figure
p = figure(title='Sine Wave Function', x_axis_label='x', y_axis_label='y')

# Add the line plot
```

```

p.line(x, y, line_width=2)

# Add grid lines
p.grid.grid_line_alpha = 0.5

# Show the plot
output_notebook()
show(p)

```

19. Using Bokeh, generate a bar chart of randomly generated categorical data, color bars based on their values, add hover tooltips to display exact values, label the axes, and set the title as 'Random Categorical Bar Chart'.

```

[60]: from bokeh.models import HoverTool
import random

# Generate random categorical data
categories = ['A', 'B', 'C', 'D', 'E']
values = [random.randint(1, 10) for _ in range(len(categories))]

# Create a DataFrame
data = {'Categories': categories, 'Values': values}
df = pd.DataFrame(data)

# Initialize the Bokeh figure
p = figure(x_range=categories, title='Random Categorical Bar Chart',
           x_axis_label='Categories', y_axis_label='Values')

# Create the bar chart
p.vbar(x='Categories', top='Values', width=0.5, color='blue', source=df)

# Add hover tooltips to display exact values
hover = HoverTool()
hover.tooltips = [('Value', '@Values')]
p.add_tools(hover)

# Label the axes
p.xaxis.major_label_orientation = 1
p.yaxis.axis_label_text_font_style = 'normal'

# Show the plot
show(p)

```

```
[ ]:
```